

Career Compass: An intelligent job search and career Development platform

BY

PRASANNA PIDIKITI

This project is submitted to the Gannon University graduate faculty in
partial fulfillment for the Master of Science degree in Computer and Information Science.
Option: **Information Technology (IT)**

Approved:

Ronny Bazan

DR RONNY BAZAN-ANTEQUERA

Advising Professor in Charge of Research
--

Gannon University
Erie, Pennsylvania 16541

December 2025

Acknowledgements

We would like to express our heartfelt thanks to Professor Ronny Bazan Antequera for his steady guidance, valuable insights, and constant encouragement throughout the development of this work. His expertise and constructive feedback greatly contributed to shaping both our understanding and the direction of our effort.

We also would like to acknowledge Gannon University for providing the necessary academic environment, technical resources, and support that is enabling us to conduct the research successfully.

We express our deep appreciation to our colleagues and friends for the discussions, suggestions, and cooperative attitude that were so helpful in advancing this project. We also thank the open-source cybersecurity communities and developers of tools and frameworks relevant to our study, as their work enabled us to analyze and design effective security solutions.

We are also very grateful to our family and friends, who have been very patient, encouraging, and steadfastly supportive. It is their encouragement that has enabled us to carry out this task confidently to its completion.

Table of Contents

Acknowledgements.....	ii
Abstract	v
List of Figures	vi
List of Tables.....	Error! Bookmark not defined.
1. Introduction	1
1.1 Overview	Error! Bookmark not defined.
1.2 Curriculum Scope.....	Error! Bookmark not defined.
1.3 Key Stakeholder Needs	1
1.4 Product Perspective.....	Error! Bookmark not defined.
1.5 Product Position Statement	2
1.6 Summary of Capabilities.....	2
1.7 Alternatives and Competition	Error! Bookmark not defined.
1.8 Project Management Plan	Error! Bookmark not defined.
1.9 References	Error! Bookmark not defined.
2. Requirements Management	Error! Bookmark not defined.
2.1 Requirements Development Perspective.....	Error! Bookmark not defined.
2.2 Use Characteristics.....	Error! Bookmark not defined.
2.2.1 User Classes and Characteristics	Error! Bookmark not defined.
2.2.2 Use-Case Model Survey	Error! Bookmark not defined.
2.2.3 User Documentation	Error! Bookmark not defined.
2.3 Feature Attributes.....	Error! Bookmark not defined.
2.3.1 Name_of_Attribute-1	Error! Bookmark not defined.
2.3.2 Name_of_Attribute-2.....	Error! Bookmark not defined.
2.3.3 Name_of_Attribute-3.....	Error! Bookmark not defined.
2.4 Key System Features.....	Error! Bookmark not defined.
2.4.1 Key System Feature 1	Error! Bookmark not defined.
2.4.2 Key System Feature 2 (<i>and so on</i>).....	Error! Bookmark not defined.
2.5 Key Design and Implementation Constraints.....	Error! Bookmark not defined.
2.5.1 Operating Environment.....	Error! Bookmark not defined.
2.6 Interface Requirements	Error! Bookmark not defined.
2.6.1 User Interfaces	Error! Bookmark not defined.
2.6.2 Hardware Interfaces.....	Error! Bookmark not defined.
2.6.3 Software Interfaces	Error! Bookmark not defined.
2.6.4 Communications Interfaces	Error! Bookmark not defined.
2.7 Nonfunctional Requirements.....	Error! Bookmark not defined.
2.7.1 Performance Requirements.....	Error! Bookmark not defined.
2.7.2 Security Requirements.....	Error! Bookmark not defined.
2.7.3 Software Quality Attributes	Error! Bookmark not defined.
2.7.4 Safety Requirements	Error! Bookmark not defined.
2.7.5 Other Requirements	Error! Bookmark not defined.
2.8 Assumptions and Dependencies.....	Error! Bookmark not defined.
3. Design	Error! Bookmark not defined.
3.1 Introduction	9
3.1.1 Goals and Objectives of Design.....	Error! Bookmark not defined.
3.1.2 Statement of Software Scope	Error! Bookmark not defined.
3.1.3 Software Context	Error! Bookmark not defined.
3.1.4 Major Design Constraints	Error! Bookmark not defined.
3.2 Data Design.....	Error! Bookmark not defined.
3.2.1 Major Internal Software Data Structure.....	Error! Bookmark not defined.
3.2.2 Global Data Structure	Error! Bookmark not defined.
3.2.3 Temporary Data Structure	Error! Bookmark not defined.

3.2.4	Database Description	Error! Bookmark not defined.
3.3	Architectural and Component-Level Design.....	Error! Bookmark not defined.
3.3.1	Program Structure	Error! Bookmark not defined.
3.3.2	Key Software Components	Error! Bookmark not defined.
3.4	User Interface Design.....	Error! Bookmark not defined.
3.4.1	Description of the User Interface	Error! Bookmark not defined.
3.4.2	Interface Design Rules.....	Error! Bookmark not defined.
3.4.3	Components Available.....	Error! Bookmark not defined.
3.4.4	User Interface Design Description.....	Error! Bookmark not defined.
3.5	Restrictions, Limitations, and Constraints	Error! Bookmark not defined.
4.	Verification and Validation	17
4.1	Test items	Error! Bookmark not defined.
5.	Conclusion.....	36
6.	Bibliography	46
Appendix A: Glossary		20
Appendix B: Use Case Analysis		20
Appendix C: Analysis Models		20
Appendix D: Design Models		20
Appendix E: Testing Log and Summary Status		20
Appendix F: Screen Captures		20
Appendix G: Project File Repository Definitions.....		20

Abstract

The Career and Salary Prediction System is an intelligent machine learning–based module developed to support students and professionals in making data-driven career and financial decisions. The system analyzes user information such as education, skills, and experience to provide personalized career recommendations and salary predictions aligned with current market trends.

Using supervised machine learning techniques, the system employs classification models for identifying suitable career paths and regression models for estimating salary ranges. Datasets containing career information, job profiles, and compensation data are preprocessed and utilized to train and evaluate predictive models. The solution is implemented using Python’s Flask framework for the web interface and SQLite for backend data management, ensuring a lightweight yet scalable architecture.

By combining predictive analytics with an intuitive user interface, the system enhances traditional career guidance methods by offering intelligent insights tailored to each user’s profile. It empowers individuals to understand their employability potential, plan skill development, and make informed decisions about their career trajectory and earning expectations. This project demonstrates the integration of artificial intelligence, data science, and web technologies to create a practical and reliable career advisory solution suitable for academic, professional, and employment platforms.

List of Figures

Figure 1. Maybe a Context Diagram **Error! Bookmark not defined.**
Figure 2. Maybe something else..... **Error! Bookmark not defined.**
Figure 3. This is another Caption **Error! Bookmark not defined.**

< A listing of the figures or graphs displayed in the document are itemized here, including the page number of the appearance of the graphic. The following format pattern should be used in the listing and should be used as the graphic’s caption. The text of the caption should be boldface, all capital letters, the caption should appear under the graphic, and the numbering should be enabled through the /References/Insert-Table-of-Figures option of MS Word. “N” refers to the page number where the figure appears.

The List of Figures for this template has been generated above.>

List of TABLES

Table 1. Maybe a Context Diagram	Error! Bookmark not defined.
--	-------------------------------------

1. Introduction

1.1 Background of the Study

The rapid evolution of artificial intelligence (AI) and data analytics has transformed how individuals plan, pursue, and advance in their careers. With the rise of digital recruitment platforms and online learning systems, the global job market has become more data-driven than ever. However, despite the availability of information, many job seekers still struggle to identify suitable career paths, align their skills with market demands, and make informed decisions regarding professional growth.

Traditional job portals and career guidance systems are often static, lacking the intelligence to personalize recommendations or forecast future career trajectories. Consequently, job seekers face challenges such as mismatched qualifications, inefficient job searches, and limited visibility into skill gaps. Employers, on the other hand, encounter inefficiencies in connecting with the right talent that matches their organizational needs.

The emergence of machine learning and intelligent data modeling provides a transformative opportunity to bridge this gap. By integrating predictive analytics, recommendation systems, and skill-based data visualization, intelligent platforms can empower users to make evidence-based career decisions.

This research introduces Career Compass, an intelligent web-based platform designed to enhance career development and job search efficiency. Built using Flask for web orchestration, SQLite for database management, and scikit-learn for machine learning predictions, the system aims to offer personalized career recommendations, salary predictions, and skill gap analysis. The integration of AI models and user-centric design makes Career Compass a comprehensive digital companion for students, professionals, and job providers.

1.2 Basic Introduction / Background of the Study

Career choice and income expectation are two of the most influential factors in an individual's professional development. Despite the abundance of job-related data available online, extracting meaningful and personalized guidance from this information remains challenging. Many individuals rely on generic job platforms that lack predictive intelligence or personalized recommendations.

The Career and Salary Prediction module within the Career Compass platform utilizes machine learning models to overcome these challenges. The career prediction component classifies users into suitable job categories based on factors such as educational background, skills, and aptitude data. The salary prediction component applies regression models to estimate income ranges corresponding to experience levels and job roles.

By combining these predictive insights, the system supports both fresh graduates and working professionals in understanding their career potential and realistic earning expectations, thus enabling informed decision-making and skill development planning.

1.3 Key Stakeholder Needs

The lack of intelligent, data-driven systems for career guidance has resulted in several challenges:

- Career recommendations are often based on subjective assessments rather than measurable data.

- Salary expectations vary widely and are influenced by unreliable or inconsistent online information.
- Users lack access to tools that combine skill analysis with predictive modeling for both career and salary insights.

This study aims to design and implement a machine learning-based career and salary prediction system that analyzes user profiles to recommend suitable career paths and predict expected salaries. By applying classification and regression algorithms to structured datasets, the system seeks to minimize guesswork and provide objective, evidence-based insights.

1.4 Nature and Significance of the Problem

The complexities of career decision-making and the uncertainty surrounding salary expectations have significant implications for both individuals and organizations. In today's dynamic job market, rapid changes in technology and evolving skill demands make it increasingly difficult for students and professionals to identify suitable career paths and realistic compensation levels. Traditional career counseling methods often rely on static data or subjective human assessments, which fail to reflect real-time market trends and personalized user profiles.

The significance of this research lies in its potential to advance intelligent career guidance systems by integrating machine learning-based predictive analytics. By applying data-driven models for career and salary prediction, the system enhances decision-making accuracy and bridges the gap between academic preparation and industry expectations. This approach empowers users to make informed choices about career direction, skill development, and income planning based on measurable data rather than intuition.

Through the development of the Career and Salary Prediction System, this research contributes to the growing field of career intelligence by combining web technologies and artificial intelligence to provide a personalized, accessible, and scalable solution for students, graduates, and working professionals.

1.5 Product Position Statement

The primary objective of this research is to design, develop, and evaluate an intelligent machine learning-based system that predicts suitable career paths and corresponding salary ranges for users based on their educational background, skills, and professional experience. By leveraging a combination of classification and regression algorithms within a Flask-based web framework, the study aims to create a platform—Career and Salary Prediction System—that provides personalized, data-driven insights for career development and financial planning.

The system is intended to assist users in identifying the most appropriate career opportunities, understanding their market value, and making informed decisions about skill enhancement. Through the integration of predictive analytics, data preprocessing, and interactive web interfaces, the proposed system seeks to minimize uncertainty in career planning and establish a reliable, intelligent approach to professional growth in an evolving job market.

1.6 Study Questions / Hypotheses

This research is guided by the need to explore the role of machine learning in improving the accuracy and personalization of career and salary predictions. The study focuses on understanding how predictive models can assist users in making data-informed career decisions.

Key research questions include:

- Can classification algorithms effectively predict suitable career paths based on user profiles such as education, experience, and skills?
- How accurately can regression models estimate salary ranges for different job categories and experience levels?
- To what extent can predictive modeling improve user satisfaction and confidence in career decision-making?
- Is the integration of career and salary prediction models within a web-based platform both scalable and practical for real-world use in educational or professional environments?

These questions aim to assess the predictive capability, usability, and impact of intelligent algorithms on career planning and employability enhancement.

1.7 Limitations of the Study

While the Career and Salary Prediction System present a promising step toward intelligent career guidance, the study recognizes several inherent limitations.

The accuracy and generalizability of predictions depend largely on the quality, diversity, and representativeness of the datasets used for model training. Variations in industry standards, regional salary differences, and evolving job market dynamics may influence the reliability of model outputs.

Additionally, the study focuses primarily on offline dataset analysis rather than real-time API integration with external platforms such as LinkedIn or Glassdoor. Computational resources and dataset constraints may also limit the exploration of deep learning models for enhanced prediction accuracy.

These factors may affect the system's scalability and adaptability to broader or rapidly changing employment landscapes. However, future work may include larger datasets, cloud integration, and dynamic data sources to overcome these limitations.

1.8 Definition of Terms

To ensure clarity and consistency throughout the study, the following key terms related to machine learning, predictive analytics, and career recommendation systems are defined:

- **Machine Learning (ML):** A field of artificial intelligence that enables computers to learn from data and make predictions without explicit programming.
- **Classification Algorithm:** A machine learning method that categorizes data into predefined groups or classes, used here for career prediction.
- **Regression Algorithm:** A statistical modeling technique used to estimate continuous outcomes such as salary based on input features.
- **Career Prediction:** The process of recommending suitable job roles or domains for a user based on their education, skills, and experience.
- **Salary Prediction:** A machine learning process that estimates potential income levels based on attributes such as qualification, experience, and role type.
- **Dataset:** A structured collection of data used to train and test predictive models, typically including variables such as job title, skills, and salary.
- **Flask Framework:** A lightweight Python-based web framework used to build web applications and APIs.

- **SQLite Database:** A self-contained, serverless database engine used for storing and managing structured data within the application.

This consistent terminology supports a precise understanding of the technical and analytical foundations of the system.

1.9 Literature Review

The foundation of this research draws upon a broad spectrum of studies in machine learning, career recommendation systems, and predictive analytics. Early frameworks such as Chapman et al. (2000) established the CRISP-DM (Cross Industry Standard Process for Data Mining) model, providing a systematic and iterative methodology for data understanding, preparation, modeling, and evaluation. This framework continues to serve as a cornerstone for modern data-driven projects, including career guidance and salary forecasting systems.

In the context of machine learning, Ng (2021) and Chollet (2021) provided comprehensive discussions on supervised learning techniques, highlighting how algorithms such as Decision Trees, Random Forests, and Gradient Boosting effectively handle structured, tabular datasets. These methods form the backbone of predictive modeling, particularly in problems involving categorical classification and numerical prediction—two key aspects of the present study.

Sharma and Reddy (2020) examined the use of Random Forest and Decision Tree classifiers in developing career guidance systems that analyze students' academic records, skills, and aptitude tests to predict suitable career paths. Their research demonstrated how ensemble methods could improve prediction reliability by combining multiple weak learners, thereby reducing overfitting and enhancing interpretability. Similarly, Kumar and Patel (2022) explored regression-based approaches for salary prediction, revealing that experience, educational qualifications, and professional specialization have significant predictive power over compensation levels.

Further studies by Ramesh and Sinha (2021) integrated machine learning and recommender systems to personalize career advice, leveraging both content-based and collaborative filtering techniques. Their work underscored the importance of hybrid models that adapt to individual profiles rather than relying solely on static rules. Alam and Singh (2022) expanded on this by proposing an intelligent career guidance platform using Random Forest and K-Nearest Neighbors (KNN), achieving improved classification accuracy through feature scaling and optimized hyperparameters.

Recent advancements in AI-powered employment systems, as discussed by Zhang and Lee (2023), have emphasized the convergence of user profiling, behavioral analytics, and predictive modeling to enhance employability insights. Such systems aim to identify skills gaps, forecast salary trajectories, and match individuals with emerging job roles. Gupta et al. (2023) highlighted the role of data visualization and explainable AI (XAI) in career prediction, advocating for transparency in model outcomes to foster user trust and acceptance.

Moreover, research by Das and Iyer (2022) explored the impact of feature engineering and imbalanced data handling using techniques such as SMOTE (Synthetic Minority Oversampling Technique) to ensure fair representation of minority classes in career datasets. Their findings align with the present study's use of SMOTE in the classification phase, ensuring balanced and equitable model learning.

In addition, Mishra and Roy (2023) introduced deep learning approaches, such as BiLSTM and ANN models, for career path recommendation, though they noted that these methods require large datasets and significant computational resources. Comparatively, tree-based ensemble methods like Random Forest

maintain a balance between interpretability, efficiency, and accuracy—qualities essential for web-based applications such as this one.

Building upon these foundational works, the Career and Salary Prediction System integrates machine learning, web technologies, and predictive analytics into a cohesive and interactive framework. By applying Random Forest algorithms for both career classification and salary regression, the system demonstrates the practical integration of academic research into an accessible, real-world application.

In conclusion, this literature review establishes the theoretical grounding and technological rationale for developing an intelligent, data-driven career advisory platform. Prior research validates the effectiveness of machine learning in modeling human career and compensation patterns, while also highlighting the need for interpretability, scalability, and user-centric design. The forthcoming chapters thus focus on system design, data preprocessing, model training, and evaluation, advancing the evolution of AI-assisted career guidance systems.

2. BACKGROUND AND REVIEW OF LITERATURE

2.1 Introduction

This chapter presents a comprehensive review of the existing literature related to intelligent career guidance systems and salary prediction models, providing critical insights into the technological, methodological, and analytical foundations that support this research. As the demand for data-driven career planning increases, the integration of artificial intelligence and machine learning in employment analytics has become a transformative approach. This section explores the evolution of career recommendation systems, predictive analytics for salary estimation, and the application of supervised learning algorithms in professional decision-making. By examining key research contributions and emerging trends, this chapter situates the Career and Salary Prediction System within the broader context of AI-based career development and highlights the advancements that inform its design and implementation.

2.2 Background Related to the Problem

The challenge of aligning individual qualifications with dynamic job market requirements has long been a concern in career development and human resource management. Traditional methods of career counseling—based on manual assessments or subjective advice—often fail to adapt to rapidly changing skill demands, leaving users with limited insight into viable career paths and salary expectations.

The evolution of data analytics and artificial intelligence has redefined this landscape by enabling systems capable of learning from large datasets and generating predictive insights. Early studies in data mining and employment analytics laid the groundwork for career and salary prediction, focusing on correlating educational qualifications, work experience, and industry trends.

Recent research has demonstrated that machine learning techniques, particularly classification and regression models, can effectively identify relationships between user attributes and professional outcomes. These developments have led to the creation of intelligent systems that assist users in making evidence-based career decisions. Understanding this progression underscores the motivation behind developing the Career and Salary Prediction System, which combines these technologies to enhance employability insights and bridge the gap between academic skills and labor market expectations.

2.3 Definition of Terms and Concepts

Scholars such as Ng (2021) and Chollet (2021) provide foundational definitions of key concepts central to the implementation of intelligent career and salary prediction systems. A clear understanding of these concepts ensures conceptual clarity and forms the theoretical framework upon which this system is built. These concepts collectively form the foundation for understanding how predictive modeling and web technologies combine to support career guidance and salary analysis in this project.

2.4 Literature Related to the Problem

Building upon the works of Sharma and Reddy (2020), Kumar and Patel (2022), and Zhang and Lee (2023), this section examines the key challenges and solutions proposed in the field of career prediction and salary estimation.

Sharma and Reddy (2020) emphasized the limitations of traditional career guidance methods, highlighting the lack of personalized and data-driven approaches. Their research demonstrated how supervised learning

algorithms such as Decision Trees and Random Forests can improve the accuracy of career recommendations by analyzing user skill sets, academic records, and aptitude data.

Kumar and Patel (2022) extended this approach to salary prediction, using regression techniques to analyze how educational qualifications, work experience, and job sector influence compensation trends. Their findings underscored the need for automated systems that provide reliable income forecasts to support informed career choices.

Zhang and Lee (2023) explored the integration of AI in employment analytics and discussed scalability challenges, model bias, and data diversity issues. Their research called for hybrid systems that combine machine learning insights with practical usability through web-based platforms.

Together, these studies highlight the necessity for an intelligent, user-centered system that merges predictive accuracy with accessibility—an objective directly aligned with the development of the Career and Salary Prediction System.

2.5 Literature Related to the Methodology

Methodological foundations for the Career and Salary Prediction System draw upon prior studies focused on the application of data science techniques in career analytics. Chapman et al. (2000) introduced the CRISP-DM framework, a standardized methodology for data mining projects encompassing data understanding, preparation, modeling, evaluation, and deployment. This framework serves as a guideline for the methodological structure of the present study.

Ng (2021) emphasized the importance of feature selection, data preprocessing, and algorithmic tuning to improve model performance in predictive systems. His work provides essential guidance on preparing datasets for classification and regression tasks.

Chollet (2021) highlighted practical aspects of machine learning workflow, including the role of hyperparameter optimization and model validation techniques. These insights inform the evaluation strategy employed in this project, which uses metrics such as accuracy for classification and R^2 for regression assessment.

Additionally, Patel and Mehta (2022) discussed the implementation of web-based ML systems using frameworks like Flask and databases such as SQLite, providing valuable insights for integrating predictive models into an interactive user interface.

The methodological synthesis of these studies guides the structured implementation of the Career and Salary Prediction System, ensuring a balance between predictive reliability, computational efficiency, and user accessibility.

2.6 Additional Insights and Contributions

Further research provides additional insights into the evolving intersection of artificial intelligence, data analytics, and career development. Studies such as Hussain and Rao (2023) examine the role of AI-driven recommendation systems in improving job matching and reducing underemployment, emphasizing how predictive models can align academic skills with real-world job requirements.

Li and Chen (2022) investigate the integration of psychological and behavioral data with traditional employment data to improve the personalization of career guidance. Their findings suggest that multi-dimensional datasets yield more accurate and human-centered predictions.

Emerging research also explores the use of natural language processing (NLP) for resume parsing and neural networks for advanced salary estimation. These innovations represent future directions for intelligent career advisory systems, extending beyond rule-based approaches toward deep learning and semantic data analysis.

Collectively, these contributions highlight a global shift toward AI-augmented career services and reinforce the timeliness and relevance of developing a scalable, web-integrated, and machine learning–driven Career and Salary Prediction System.

2.7 Summary

In summary, this chapter has provided an extensive review of the theoretical and practical foundations that inform the Career and Salary Prediction System. It traces the evolution of career analytics from traditional counseling methods to data-driven, machine learning–based solutions, emphasizing the growing importance of predictive modeling in employment guidance.

The chapter began with an exploration of the background related to the problem, identifying key limitations of existing career advisory systems. It then defined critical concepts such as classification, regression, and data preprocessing, establishing the conceptual framework for the system.

Subsequent sections reviewed key studies addressing challenges in prediction accuracy, model interpretability, and user accessibility. The methodological discussion connected foundational data science principles with practical implementation techniques in web-based environments. Additional insights from recent literature highlighted emerging trends such as NLP integration, behavioral data modeling, and AI-driven recommendation systems.

Together, these perspectives form a cohesive academic foundation for the design and development of the Career and Salary Prediction System, supporting its goal of enhancing employability, career clarity, and salary transparency through the power of intelligent data analytics.

3. METHODOLOGY

3.1 Introduction

In this pivotal chapter, we explore the research design, development methodology, and technical tools employed to construct the Career and Salary Prediction System, an intelligent platform that leverages machine learning to assist users in making data-driven career and financial decisions. This section presents a comprehensive overview of the applied research strategy, combining principles from data science, predictive analytics, and web application development. The methodology is closely aligned with the study's objectives—to design and evaluate a system capable of predicting suitable career paths and salary ranges based on user-specific attributes such as education, skills, and experience.

To ensure the effectiveness and feasibility of the system, a design-based research approach was adopted. This iterative process involved data preprocessing, model training, validation, and integration within a web-based platform. The practical implementation is demonstrated through a Flask-based web application that allows users to input their details, receive predictive recommendations, and visualize analytical insights. This methodological framework ensures the system's practical relevance and provides a strong foundation for scalability, performance assessment, and future enhancement.

3.2 Research Design of the Study

The research design for the Career and Salary Prediction System follows a development-oriented and experimental approach, integrating supervised machine learning algorithms with web technologies to achieve accurate prediction and user accessibility. The design process emphasizes model reliability, data accuracy, and seamless user interaction.

Central to this design is the dual-model framework:

Career Prediction Model — employs classification algorithms such as Random Forest and Decision Tree to predict suitable career paths based on user data (skills, education level, and aptitude factors).

Salary Prediction Model — utilizes regression algorithms such as Linear Regression and Random Forest Regressor to estimate salary ranges based on job role, experience, and educational background.

The models are trained and validated using structured datasets that include fields such as education, skills, experience level, and current job role. Data preprocessing steps, including normalization, encoding of categorical variables, and feature selection, ensure consistency and quality across all model inputs.

The system's user interface is built using the Flask web framework, which facilitates interaction between the user and the backend machine learning models. Data storage and management are handled using the SQLite database, ensuring lightweight and efficient local persistence.

This integrated design allows the Career and Salary Prediction System to operate as a unified platform for personalized career guidance, effectively demonstrating how AI-driven analytics can enhance employability and career planning in real-world contexts.

3.3 Data Collection Techniques and Procedures

The data collection phase of this study focuses on gathering relevant datasets and transforming them into a usable form for machine learning model training and testing. Unlike traditional surveys, this research utilizes open-source datasets containing career attributes, salary information, and skill parameters.

Data were collected from publicly available repositories such as Kaggle, Glassdoor datasets, and academic sources containing structured records of job profiles, experience levels, and salary distributions. Each dataset was evaluated for completeness, consistency, and relevance to the study's predictive objectives.

The data collection process involved the following steps:

Dataset selection and cleaning: Removal of duplicates, missing values, and irrelevant attributes.

Data transformation: Conversion of categorical variables (e.g., education, role) into numerical format using encoding techniques.

Feature extraction: Identification of critical variables influencing career and salary outcomes.

Dataset splitting: Division of data into training and testing sets to ensure model validation and prevent overfitting.

Additional synthetic data were generated for testing various user scenarios to simulate real-world conditions, such as varying skill levels and experience ranges.

Performance metrics (accuracy, precision, recall, and R^2) were computed after model evaluation to assess effectiveness.

All collected and processed data were securely stored in the project's local environment and integrated with the SQLite backend for model deployment within the web application.

3.4 Sampling Techniques / Target Population / Sample Size

The sampling process for the Career and Salary Prediction System is based on a purposeful and representative sampling technique rather than random sampling, due to the nature of the data-driven approach. The target population comprises students, recent graduates, and professionals from various academic and technical disciplines seeking career recommendations or salary insights.

The selected datasets represent multiple industries such as technology, finance, healthcare, and education, covering a broad range of experience levels and salary brackets. This diversity ensures that the trained models generalize well across different job markets and user profiles.

The sampling design emphasizes:

Inclusion of multiple job roles and industries to improve model robustness.

Representation of various education levels and experience categories.

Use of both training and testing samples (typically in an 80:20 ratio) to validate model accuracy.

This approach ensures that the models are evaluated under diverse conditions, improving the reliability and general applicability of the Career and Salary Prediction System for real-world users.

3.5 Instrument Development (Survey Questionnaire / Interviews)

Unlike traditional social science research that relies on survey questionnaires or interviews, this study utilizes computational instruments—datasets, trained algorithms, and system interaction logs—as the primary tools

for evaluation. The instruments are designed to assess model performance, prediction accuracy, and user experience in an objective and reproducible manner.

The development of these instruments involved:

Machine Learning Models: Classification and regression algorithms trained on structured datasets to generate predictive outputs.

Performance Metrics: Evaluation tools such as accuracy, precision, recall, mean absolute error (MAE), and R^2 score.

User Interface Testing: Simulation of multiple user profiles and input scenarios to verify prediction consistency and responsiveness.

Data Logging Mechanisms: Flask-based backend logging of user requests, model responses, and computation times for performance analysis.

These instruments collectively ensure a scientific and unbiased evaluation process. By relying on algorithmic validation and data-driven testing, the research achieves a higher degree of precision and repeatability than conventional qualitative assessment methods.

3.6 Instrument Reliability / Validity

The reliability and validity of the instruments used in this study were ensured through multiple validation stages.

Reliability was established by conducting repeated model training and testing cycles using different random seeds and dataset partitions. Consistent performance metrics across runs indicated model stability. Cross-validation techniques were applied to verify that the models produced reliable predictions across diverse data subsets.

Validity was assessed in two key forms:

Predictive Validity: Confirming that the classification and regression models correctly predict career paths and salary ranges based on real or test data.

Construct Validity: Ensuring that the input variables (skills, experience, education) logically and statistically relate to the predicted outcomes.

Performance measures such as confusion matrices, error analysis, and cross-validation scores were used to validate model robustness.

Together, these evaluation procedures confirm that the Career and Salary Prediction System produces consistent, credible, and generalizable results.

3.7 Implementation

The implementation phase marks the transformation of the proposed Career and Salary Prediction System from a conceptual design into a fully functional software solution. This stage involves developing machine learning models, integrating them into a web-based application, establishing database connectivity, and ensuring smooth interaction between all system components. The implementation follows a layered architecture to maintain modularity, scalability, and ease of future enhancements.

The project is implemented using the following layered structure:

1. Machine Learning Models (Python & scikit-learn)

The backbone of the system lies in the machine learning models that perform both career prediction (classification) and salary prediction (regression).

Data Preprocessing

To prepare the dataset for accurate model training:

- **Data Cleaning:** Removal of missing values, duplicates, and inconsistencies.
- **Feature Encoding:** Categorical attributes such as education level, job domain, or skills are converted using Label Encoding and One-Hot Encoding.
- **Normalization/Standardization:** Numerical features like experience and age are scaled to ensure equal contribution during model training.
- **Feature Engineering:** Derived parameters such as skill count, domain relevance score, or performance index may be added to improve model performance.

This preprocessing is implemented using pandas, NumPy, and scikit-learn preprocessing utilities.

Career Prediction (Classification Models)

The classification module identifies the most suitable job domain for a user based on their profile.

Random Forest Classifier:

Aggregates multiple decision trees to reduce overfitting.

Provides high accuracy and handles non-linear feature relationships.

Decision Tree Classifier:

Offers interpretable decision paths.

Used for comparison and fallback prediction.

Hyperparameter tuning is performed using GridSearchCV to optimize model depth, number of estimators, and feature splits.

Salary Prediction (Regression Models)

Salary ranges are predicted based on the classified domain and user-specific attributes.

- **Linear Regression:**

Captures linear trends between experience, skills, education, and salary.

- **Random Forest Regressor:**

Handles complex non-linear salary variations across job roles.

Performance metrics such as MAE, MSE, RMSE, and R^2 score are calculated to validate model quality.

Model Serialization

After training, models are saved using joblib for quick loading during runtime.

This enables the Flask application to load models instantly without re-training, improving performance and reducing server overhead.

2. Web Application Layer (Flask Framework)

The machine learning models are integrated into a Flask-based web application that serves as the main user interaction platform.

Backend (Flask)

Flask routes handle:

- Form submissions (user profile inputs)
- Model prediction calls
- Displaying prediction results
- Logging activities into the database
- Input validation ensures that all required fields are present before predictions are executed.
- Error handlers are implemented to manage invalid inputs, model failures, or database exceptions.

Frontend (HTML, CSS, Bootstrap)

The user interface is designed to be clean, responsive, and easy to navigate.

- HTML templates: Structure pages for input, output, and user interactions.
- Bootstrap components: Provide mobile-friendly layouts, cards, buttons, and form design.
- CSS styling: Adds branding, theme consistency, and visual clarity.

The frontend communicates with Flask via POST requests, ensuring seamless interaction with backend logic.

3. Database Layer (SQLite & SQLAlchemy)

The system incorporates persistent storage to enhance usability and maintain analytical data.

SQLite Database

SQLite is chosen due to:

- Lightweight deployment
- Server-less architecture
- Support for structured relational data

Data Stored Includes:

- User registration details
- Prediction inputs and outputs
- Login sessions
- Feedback logs
- System metrics for future model refinement

SQLAlchemy ORM

SQLAlchemy is used for:

- Table creation via Python models
- Secure CRUD operations
- Mapping Python objects directly to database entries
- Abstracting raw SQL queries for maintainability

This ensures smoother database management and enhances portability.

4. Integration and Deployment

Model Integration

- Trained and serialized models are loaded within Flask startup scripts.
- Prediction logic is embedded inside route functions.
- Preprocessing steps (encoding, scaling) are replicated during prediction for consistency with training.

Testing

Comprehensive testing is performed:

- **Functional Testing:**
Ensures all user inputs, navigation, and prediction functionalities work correctly.
- **Model Validation:**
Confirms prediction accuracy on unseen test data.
- **Interface Testing:**
Checks responsiveness across devices and browsers.

Deployment Setup

The system is deployed locally during development:

- Flask servers run in debug mode for rapid development.
- Static resources (CSS, JS, images) organized into separate directories.
- Optionally configured for deployment on cloud platforms (Heroku, AWS EC2) using environment-specific configurations.

5. System Characteristics

The layered design delivers:

- **Flexibility:** Easy to replace or retrain ML models.
- **Scalability:** New prediction modules or dashboards can be added.
- **Maintainability:** Clear separation of concerns across layers.
- **Performance:** Fast real-time predictions due to model serialization and efficient routing.

3.8 Confidentiality and Ethics / Ethical Approval

The research adheres to strict ethical guidelines concerning data usage, privacy, and integrity. All data sets used in this study are open-source and anonymized, containing no personally identifiable information. The project fully complies with academic data ethics standards by maintaining transparency in data handling and ensuring that no private user data is collected or stored.

Ethical principles implemented include:

Data Anonymity: Use of de-identified datasets to ensure user privacy.

Data Integrity: Ensuring accuracy and reliability through verification and preprocessing.

Transparency: Documenting all algorithms, preprocessing steps, and validation methods for reproducibility.

Formal ethical approval was not required, as the study did not involve direct human participants. However, ethical standards were upheld throughout the research to maintain fairness, accountability, and respect for data integrity.

3.9 Pilot Study

The pilot phase of the Career and Salary Prediction System focused on testing the integrated workflow from data input to output prediction within a controlled environment. The objective was to validate the functionality, user experience, and accuracy of predictive components before full-scale deployment.

The pilot implementation included:

Testing multiple user input scenarios (varied education, skills, experience).

Verifying the correctness of career and salary predictions.

Assessing system responsiveness and error handling.

Ensuring the interface provides intuitive guidance and real-time feedback.

Feedback from the pilot phase was used to refine data preprocessing, optimize model performance, and enhance interface usability. This stage confirmed the system's readiness for broader deployment and validated its ability to assist users in informed career planning.




Figure 3.1 System Architecture

3.10 Summary

This chapter detailed the methodological framework adopted for the design, development, and evaluation of the Career and Salary Prediction System. The research follows a development-based approach integrating machine learning, data preprocessing, and web development techniques.

Data collection was conducted through curated and cleaned datasets, while data-driven instruments—rather than surveys—were used to evaluate model and system performance. Reliability and validity were confirmed through multiple testing cycles and cross-validation.

The implementation process combined machine learning models, Flask-based backend logic, and SQLite database management to build an interactive web platform. Ethical considerations ensured responsible data handling, and the pilot study provided valuable insights into usability and functionality.

Collectively, this methodology establishes a solid foundation for developing intelligent, data-driven career guidance tools that enhance employability and enable evidence-based decision-making.

4. DATA PRESENTATION AND ANALYSIS

4.1 Introduction

This chapter presents the empirical results obtained after training, testing, and integrating the Career & Salary Prediction System into the application workflow. The analysis focuses on two pipelines that both use Random Forest:

Career classification, which predicts the most suitable career from user attributes, and

Salary regression, which estimates expected salary from demographic and job-related features.

Evidence is drawn from the system's Python workflows—preprocessing logs, SMOTE balancing on the training fold (for classification), model evaluation outputs, and the visualizations generated by the notebooks/scripts. The figures referenced in this chapter are:




Figure 4.1 – Confusion Matrix (Career RF)




Figure 4.2 – Feature Importance (Career RF)




Figure 4.3 – Classification Metrics Bar (Career RF)



Figure 4.4 – Predicted vs. Actual (Salary RF)




Figure 4.5 – Salary Residuals Histogram (Salary RF)

Together, these outputs demonstrate how the proposed pipelines behave with real data and how the models generalize to unseen examples.

4.2 Data Collection

Two curated datasets were used:

Career Path Dataset (corrected_career_path.csv): contains a categorical Field attribute and multiple numeric/ordinal attributes that summarize aptitude, academic indicators (e.g., GPA), projects, internships, certifications, and soft-skill scores. The target label is Career.

Salary Dataset (Salary Data.csv): contains Age, Gender, Education Level, Job Title, Years of Experience, and the continuous target Salary.

Both datasets were inspected for completeness and consistency. Rows missing the target variable (i.e., Career or Salary) were removed. Categorical columns were normalized to string type for encoding; numeric columns were checked for outliers and missingness and then standardized for model consumption.

4.3 Sampling Technique and Inclusion/Exclusion Criteria

A purposeful, representative sampling strategy was used to ensure fair coverage across target classes and feature ranges.

Inclusion criteria (Career): records with valid Career labels and complete entries for all required predictors.

Inclusion criteria (Salary): records with non-missing Salary and all required predictors present.

Exclusion criteria: rows missing the target, corrupted records, or entries with irreparable categorical/format issues.

For the career task, a stratified train–test split preserves class proportions. SMOTE is applied only to the training set to alleviate class imbalance and avoid data leakage. For the salary task, a standard train–test split is used.

4.4 Data Importing, Handling, and Validation

CSV files are loaded directly by the Python scripts. The following preprocessing is applied:

Categorical encoding with LabelEncoder (e.g., Field, Career, Gender, Education Level, Job Title).

Numeric scaling with StandardScaler for continuous predictors (e.g., Age, Years of Experience, GPA and skill/aptitude scores).

Imbalance handling (classification): SMOTE on the training fold only.

Input validation: data types are coerced; empty strings and NaNs are handled; encoders and the scaler are persisted to ensure consistency at inference time.

This sequence guarantees that the model sees a clean, consistent representation and that the deployed Flask routes can reproduce the same preprocessing during online predictions.

4.5 Questionnaire Limitation

No qualitative user surveys/interviews were conducted for this iteration. Evaluation relies on objective telemetry—confusion structure, aggregate metrics, feature importance, predicted-vs-actual behavior, and error distribution. Future versions may include user studies for interpretability and UI feedback.

4.6 Implementation Summary

Career pipeline: encode → scale → SMOTE (train) → Random Forest Classifier → metrics, confusion matrix, and feature importance → persist (best_career_model.joblib, encoders, scaler).

Salary pipeline: encode/scale → Random Forest Regressor → predicted-vs-actual and residual plots → persist (salary_predictor.pkl bundle with model, encoders, scaler, and metrics).

All figures in this chapter are exported from the same scripts to ensure reproducibility.

4.7 Alternative Algorithms Consideration

Although this study employs Random Forest as the exclusive machine learning algorithm for both career classification and salary prediction, several alternative approaches were initially reviewed to evaluate their suitability for the system's requirements. These alternatives included Decision Tree, Logistic Regression, Support Vector Machine (SVM), and Gradient Boosting methods. Each algorithm offers distinct advantages and trade-offs in terms of interpretability, scalability, and computational efficiency.

4.7.1 Decision Tree

The Decision Tree algorithm was initially considered due to its simplicity, transparency, and ability to generate rule-based decisions that can be easily visualized and interpreted by non-technical stakeholders. Decision Trees construct a hierarchical structure of decisions, providing an intuitive understanding of how each input feature contributes to the outcome. However, despite their interpretability, Decision Trees are highly prone to overfitting, particularly in high-dimensional datasets with both categorical and numerical features. Preliminary experiments conducted during the study revealed that Decision Trees achieved lower predictive accuracy and exhibited high variance on unseen data compared to ensemble-based approaches. Consequently, while Decision Trees offered a clear baseline for understanding feature relationships, they were deemed unsuitable as the final model due to their limited generalization capacity.

4.7.2 Logistic Regression

Logistic Regression was evaluated as a baseline model for the career classification task owing to its mathematical simplicity, fast computation, and well-established interpretability. It models the probability of class membership as a logistic function of input features, making it effective for problems with linearly separable classes. However, the model assumes a linear relationship between predictors and the outcome variable, which restricts its ability to capture complex non-linear interactions between academic, technical, and behavioral attributes such as aptitude scores, GPA, certifications, and communication skills. During preliminary testing, Logistic Regression produced lower accuracy and recall rates compared to Random Forest, particularly for multi-class outputs. Thus, while suitable as a reference model for benchmarking, Logistic Regression was not sufficient to represent the intricate feature dependencies present in the dataset.

4.7.3 Support Vector Machine (SVM)

The Support Vector Machine (SVM) algorithm was also considered for its solid theoretical foundations and strong performance in high-dimensional feature spaces. SVMs utilize kernel functions to handle non-linear relationships, making them effective for classification tasks involving complex decision boundaries. However, they are computationally intensive, especially for large datasets, and require careful tuning of kernel parameters such as C and gamma to achieve optimal performance. Furthermore, extending SVMs to

multi-class classification (as required in the career prediction task) significantly increases both model complexity and training time. Additionally, the interpretability of SVMs is relatively low compared to tree-based models, making them less practical for user-facing predictive systems where model transparency is crucial. For these reasons, SVMs were not selected as the primary algorithm.

4.7.4 Gradient Boosting (XGBoost / AdaBoost)

Lastly, Gradient Boosting algorithms such as XGBoost and AdaBoost were explored due to their proven success in structured and tabular data prediction tasks. These ensemble methods build successive weak learners (typically Decision Trees) in a stage-wise manner to reduce model error iteratively, often achieving high predictive accuracy. However, these models are computationally demanding, sensitive to hyperparameters, and can overfit easily without careful regularization and parameter tuning. Moreover, they require more extensive training time and computational resources than Random Forest, making them less ideal for deployment within a lightweight Flask-based web framework. Given the project's emphasis on robustness, scalability, and ease of integration, Random Forest was preferred over Gradient Boosting approaches, offering a stable and efficient alternative with minimal hyperparameter sensitivity.

4.7.5 Rationale for Selecting Random Forest

After evaluating the alternative algorithms, **Random Forest** was ultimately selected as the most suitable model for both tasks—**career classification** and **salary regression**—due to its proven ability to deliver a strong balance between predictive accuracy, interpretability, and computational efficiency. The decision was guided by both empirical results from pilot experiments and theoretical considerations of scalability and robustness in real-world data environments.

- **Robustness to Noise and Outliers:**

Random Forest combines the results of multiple decision trees trained on random subsets of data and features, effectively minimizing the influence of noisy, unbalanced, or outlier-prone samples. This ensemble averaging mechanism enhances model stability and ensures consistent performance across diverse input distributions.

- **Ability to Handle Mixed Data Types:**

Unlike many algorithms that require extensive preprocessing or feature transformation, Random Forest inherently supports both categorical and continuous variables. This flexibility simplifies the data preparation process, allowing seamless integration of encoded features such as education level, job title, gender, and quantitative indicators like GPA or years of experience.

- **Resistance to Overfitting:**

By employing **bootstrap aggregation (bagging)** and random feature selection at each node split, Random Forest mitigates the high variance typically associated with individual decision trees. This mechanism ensures better generalization to unseen data, yielding reliable predictions for both career paths and salary estimations in practical use.

- **Interpretability and Feature Insights:**

A key advantage of Random Forest lies in its ability to compute **feature importance scores**, which quantify the relative influence of each predictor variable. This interpretability provides valuable transparency, enabling users and analysts to understand which academic, experiential, or skill-based factors most strongly shape model outputs—an essential quality for decision-support systems in education and career counseling.

4.8 Data Presentation – Features and Variables

The data sets used for the Career and Salary Prediction System were carefully selected and structured to capture both academic and professional dimensions relevant to predicting career paths and expected salaries. Each dataset underwent preprocessing steps such as categorical encoding, numeric scaling, and consistency verification to ensure accurate model training and deployment readiness.

For the career classification model, the dataset (`corrected_career_path.csv`) includes a diverse collection of attributes representing academic background, technical proficiency, and interpersonal competencies. The key predictors comprise the Field of study, which defines academic discipline such as Engineering, Science, Commerce, or Arts and serves as one of the most dominant features influencing career recommendations. Alongside this, the Grade Point Average (GPA) provides a numeric measure of academic performance, reflecting intellectual aptitude and learning consistency. Features such as Projects and Internships quantify practical exposure and real-world experience, which are crucial indicators of employability and readiness for professional roles.

Additional predictors include Certifications, representing specialized credentials like programming or domain-specific training (e.g., cloud computing, cybersecurity), which strengthen job alignment for technical careers. The model also incorporates several skill-based attributes—including Coding Skills, Problem-Solving Skills, and Analytical Skills—that represent a candidate's cognitive and logical reasoning capabilities. Equally significant are the soft skill indicators, such as Communication, Presentation, Teamwork, and Leadership Qualities, which collectively describe interpersonal and managerial effectiveness. Features like Networking Skills, Extracurricular Activities, and Industry Certifications capture holistic personality traits, advanced learning, and research involvement, contributing additional predictive depth.

The target variable for this dataset is Career, which represents the most suitable professional role for the candidate—examples include Software Engineer, Data Analyst, Civil Engineer, HR Specialist, Graphic Designer, or Professor. This variable is label-encoded during model training and then inverse-transformed into its original text form when generating predictions. The diverse blend of quantitative (e.g., GPA, project count) and qualitative (e.g., communication and leadership ratings) variables allow the Random Forest Classifier to detect multi-dimensional relationships between academic achievement, skills, and career outcomes.

The second dataset, used for the salary regression model (`Salary Data.csv`), focuses on demographic and job-related features that collectively determine expected salary levels. Among the most important predictors is Age, representing overall career maturity and indirectly correlating with years of professional experience. Gender is included as a demographic variable, primarily for analytical completeness rather than decision bias, allowing the model to capture distributional patterns in salary structures across different demographics. The Education Level feature records the highest academic qualification attained by the individuals such as Bachelor's, Master's, or Ph.D.—which often serves as a strong determinant of compensation potential, especially in technical and managerial roles.

Another critical predictor is the Job Title, which identifies the professional role or designation, for instance, Software Developer, Data Scientist, Project Manager, or Marketing Executive. It directly affects salary predictions because compensation levels vary widely across roles and industries. Finally, Years of Experience quantifies professional exposure and acts as one of the most influential factors for salary estimation, as experience tends to correspond strongly with pay increments and seniority levels.

In both datasets, categorical variables such as Field, Gender, Education Level, and Job Title were encoded numerically using Label Encoder, while continuous variables such as GPA, Age, and Years of Experience were normalized with StandardAero to maintain uniformity during model training. These transformations ensure that features of different scales and formats contribute fairly to the learning process. By combining academic, skill-based, and experiential attributes, the feature sets for both career classification and salary regression provide a comprehensive foundation for reliable and interpretable prediction using the Random Forest algorithm.

4.9 Rationale for Feature Selection

The selection of features for both the career classification and salary prediction models was guided by empirical research, HR analytics best practices, and practical considerations of data availability, interpretability, and predictive value. Each feature was chosen to capture critical dimensions of employability and compensation potential, namely, academic background, technical competence, professional experience, and interpersonal aptitude all of which collectively shape an individual's career trajectory and salary prospects.

Prior studies and industry insights consistently identify education level, years of experience, and job role as primary determinants of salary outcomes, reflecting the traditional economic relationship between human capital accumulation and market valuation. Accordingly, these features were incorporated into the Salary Regression model, alongside demographic attributes such as age and gender, which—while ethically neutralized during analysis—can provide context to existing real-world disparities in pay and opportunity structures. The inclusion of education level ensures that differences in academic qualifications (e.g., Bachelor's, Master's, Doctorate) are appropriately captured as stratifying factors in compensation modeling, while years of experience encapsulates learning curves, productivity growth, and career seniority effects.

For the Career Classification model, the chosen predictors—field of study, GPA, certifications, projects, internships, and aptitude/skill-based indicators (such as coding, analytical thinking, communication, and teamwork) reflect a multidimensional understanding of employability. Field of study anchors the candidate's academic orientation, while GPA and certifications provide quantifiable evidence of domain mastery and continuous learning. The inclusion of projects and internships serves as a proxy for applied experience and professional readiness, bridging the gap between theoretical education and workplace performance. Furthermore, soft-skill measures like teamwork, leadership, and communication scores contribute to capturing the non-technical attributes increasingly emphasized in modern recruitment analytics.

These features were not arbitrarily chosen; rather, they were validated through iterative model testing, feature-importance evaluation, and cross-referencing with existing literature in workforce analytics and predictive hiring. Random Forest's inherent ability to rank feature significance further supported the interpretive justification for their inclusion, confirming that education, experience, and role-related skills consistently exhibited high information gain across both predictive tasks.

Finally, the chosen feature set was intentionally designed for user accessibility and web integration. Each variable corresponds to inputs that can be easily provided by users through the system's web interface, such as education level, job title, or years of experience. This ensures not only data relevance but also practical usability, allowing real-time interaction with the trained models without requiring complex or sensitive personal information.

4.10 Algorithms

The Career and Salary Prediction System rely on the **Random Forest algorithm** as the primary machine learning technique for both classification (career recommendation) and regression (salary prediction). Random Forest is a highly flexible ensemble method capable of modeling complex patterns in structured,

mixed-type datasets. Its ability to handle high-dimensional, noisy, and heterogeneous features makes it an ideal choice for this project, where academic records, demographic details, skills, and professional attributes all interplay to shape career and salary outcomes.

4.10.1 Overview of the Random Forest Algorithm

Random Forest is an **ensemble-based learning technique** that constructs a collection of independent decision trees during training and combines their outputs to produce a final prediction. The underlying principles include:

- **Bootstrap Aggregation (Bagging):**
Each tree is trained using a random subset of the training dataset drawn with replacement. This improves variance reduction and model stability.
- **Random Feature Selection:**
At each tree split, a random subset of features is considered rather than using all features. This decorrelates the trees and improves generalization.
- **Aggregation of Predictions:**
 - For **classification**, the final output is determined via **majority voting**.
 - For **regression**, the final prediction is the **average** of all three outputs.

These mechanisms help reduce overfitting, improve predictive accuracy, and increase robustness to noise.

4.10.2 Algorithm Workflow in the System

Dataset Preparation

- Input features include:
academic field, GPA/percentage, skill ratings, certifications, technical skills, soft skills, demographic details, and behavioral parameters.
- Data preprocessing ensures:
 - Categorical features are label-encoded.
 - Numerical features are normalized where required.
 - Missing values are imputed.
 - Data is split into training and testing subsets for evaluation.

Tree Construction

Each decision tree within the Random Forest is constructed using:

- **Gini Impurity or Entropy (for classification)** to decide optimal splits.
- **Mean Squared Error (for regression trees)** to minimize prediction error.
- **Depth constraints** to avoid overly complex trees.
- **Random feature subsets** at each split to ensure diversity across the ensemble.

Model Training

Training is performed using scikit-learn's:

- RandomForestClassifier

- RandomForestRegressor

The models learn patterns from historical data using:

- Bootstrapped sampling
- Recursive binary splitting
- Aggregated decision-making

Hyperparameter tuning employs methods such as:

- Grid Search
- Randomized Search
- Cross Validation (K-Fold)

Key hyperparameters tuned include:

- **n_estimators**: Number of trees
- **max_depth**: Depth of each tree
- **max_features**: Number of features considered at each split
- **min_samples_split** / **min_samples_leaf**: Controls tree growth
- **bootstrap**: Whether sampling is with replacement

4.10.3 Career Prediction: Random Forest Classifier

The career prediction module uses a Random Forest Classifier to recommend suitable job roles or domains based on user attributes.

Why It Works Well for Career Prediction

- Handles mixed data types (numeric + categorical).
- Captures non-linear interactions between skills and academic features.
- Reduces overfitting compared to a single decision tree.
- Identifies **feature importance**, enabling interpretability.

Important Features Identified (Example Interpretation)

Random Forest often highlights top contributors such as:

- Technical skill ratings (e.g., programming, data analysis)
- Soft skills (communication, teamwork, problem-solving)
- Academic performance
- Certifications / domain knowledge
- Practical experience (internships, projects)

This allows the system to not only predict careers but also explain *why* a particular domain is recommended.

Prediction Mechanism

1. Each tree independently predicts a career domain.
2. The final output is the category with the highest vote.
3. Confidence scores may be computed based on vote proportions.

4.10.4 Salary Prediction: Random Forest Regressor

The salary estimation module uses Random Forest Regression to model salary as a continuous numeric prediction.

Why It Works Well for Salary Prediction

- Captures complex salary patterns across industries.
- Resistant to outliers (e.g., unusually high salaries).
- Models' non-linear relationships between features like experience, education level, job role, age, location, etc.
- Requires minimal data preprocessing.

Regression Behavior

- Each decision tree outputs a salary estimate.
- The model averages these values for a final stable prediction.
- Reduces noise and variation inherent in salary data.

4.10.5 Model Evaluation Techniques

Classification Metrics

To assess career prediction effectiveness:

- **Accuracy:** Measures of overall correctness.
- **Precision (Weighted):** Ensures class-specific reliability.
- **Recall (Weighted):** Ensures minority classes are not ignored.
- **F1-score (Weighted):** Balances precision and recall.

Visual Diagnostics

- **Confusion Matrix:** Identifies misclassifications and domain overlaps.
- **Feature Importance Graphs:** Reveal which factors influence career outcomes.
- **ROC Curves (optional):** Evaluate class-level separability.

Regression Metrics

To evaluate salary prediction accuracy:

- **RMSE (Root Mean Squared Error):**
Lower RMSE indicates closer salary predictions to actual values.
- **R² (Coefficient of Determination):**
Measures how much salary variation the model explains.
- **Visual Diagnostics**
 - **Actual vs. Predicted Scatter Plot:**
Shows how close predictions are to true salaries.
 - **Residual Plot:**
Helps detect biases such as skew or heteroscedasticity.
 - **Error Distribution Curve:**
Ensures no systematic prediction errors exist.

4.10.6 Advantages of Using Random Forest in This Project

- High accuracy for both classification and regression.
- Robust to noisy, missing, and imbalanced data.
- Supports large feature sets without dimensionality reduction.
- Provides interpretability via feature importance.
- Requires minimal parameter tuning to achieve good performance.
- Suitable for real-time predictions in the Flask application.

4.10.7 Limitations and Considerations

Although Random Forest performs exceptionally well, it has some constraints:

- Large models can be slower during prediction, though acceptable for this system.
- Not ideal if extremely high interpretability is required (compared to single decision trees).
- May require significant memory for large datasets due to many trees.

These limitations are mitigated through controlled hyperparameter tuning and model optimization.

.

4.11 Visualizations and Results

The visual evaluation of the Career and Salary Prediction System provides critical insights into model behavior, strengths, and limitations. The following graphical analyses were generated during experimentation with the Random Forest classifier and regressor. Each visualization contributes to understanding how well

the model generalizes, how predictions relate to actual patterns in the dataset, and where further refinement is needed.

4.11.1 Confusion Matrix: Random Forest (Career)

- **Detailed Interpretation**

The confusion matrix displays the distribution of correct and incorrect predictions across all career classes. A strong diagonal line—where predicted classes match actual classes—indicates that the model achieves high classification accuracy across most categories.

Key observations:

- **Strong Diagonal Structure:**
Most predictions fall along the diagonal, confirming that the Random Forest classifier identifies appropriate career domains with high reliability.
- **Localized Off-Diagonal Patterns:**
Misclassifications tend to occur between roles that share overlapping skillsets or academic profiles. For example:
 - **Software Developer ↔ Data Analyst**
 - **Network Engineer ↔ System Administrator**
 - **Data Scientist ↔ Machine Learning Engineer**
These overlaps reflect realistic career adjacency rather than model weakness.

- **Meaningful Confusion Zones:**

The confusions highlight areas where:

- Users possess hybrid skill profiles.
- Certain job categories have overlapping competencies (e.g., Python, SQL, cloud skills).
- Data representation is uneven (minority roles may be underrepresented).

- **Opportunities for Enhancement:**

Adding more detailed features—such as project complexity, specific tools (TensorFlow, AutoCAD, Power BI), internship duration, or domain certifications—would sharpen class boundaries.

Overall, the matrix reinforces that the model is effective and logically consistent with real-world career similarity patterns.

4.11.2 Feature Importance: Random Forest (Career Prediction)

- **Detailed Interpretation**

The feature importance of visualization reveals how much each input feature contributed to the classification decisions. Random Forest computes importance based on impurity reduction across all trees.

Key insights:

- **Field of Study Dominates Predictive Power:**
Academic background naturally aligns with early-stage career pathways (e.g., CSE → Software roles, ECE → Embedded/Networking roles).
Its high importance validates the domain knowledge built into hiring trends.

- **GPA and Academic Performance Indicators:**
Moderately high importance values for GPA or CGPA indicate that academic consistency still plays a meaningful role especially for structured career tracks such as analytics, engineering, and technical consulting.
- **Projects, Internships, and Technical Courses:**
These features contribute significantly, suggesting the model values hands-on experience and specialization.
Students with relevant project portfolios receive clearer predictions.
- **Soft Skills (Teamwork, Communication, Problem-solving):**
Features related to soft skills also show non-trivial influence. This proves that Random Forest is not purely relying on technical metrics but captures holistic candidate profiles.
- **Distributed Importance:**

No single feature overwhelmingly dominates, confirming:

- The model is robust
- Each attribute contributes meaningfully
- Career prediction requires **multi-dimensional** evaluation

This distribution indicates a healthy balance and aligns with actual hiring criteria for entry-level positions.

4.11.3 Random Forest Career Classification Metrics

- **Detailed Interpretation**
-

Evaluation metrics Accuracy, Precision, Recall, and F1-Score—provide a statistical summary of model performance.

Key observations:

- **High Accuracy Across Classes:**
The classifier consistently predicts the correct career category for most users.
- **Balanced Precision and Recall:**

The narrow gap between these metrics implies:

- The model does not overpredict popular classes.
- Minority career categories (e.g., UI/UX, HR Technical Specialist, Cybersecurity Analyst) also receive sufficient sensitivity.

- **Strong F1-Score:**

The F1-score reflects the harmonic meaning of precision and recall. Its high value suggests the model nicely balances performance across uneven class distributions.

- **Impact of SMOTE:**

The improvement in minority-class representation is visible. SMOTE (Synthetic Minority Oversampling Technique) reduces bias toward majority career categories, enabling more equitable performance.

- **No Evidence of Severe Overfitting:**

Model performance remains stable across training and testing datasets.

Collectively, these metrics validate that the career classifier is reliable, generalizes well, and treats all career classes fairly.

4.11.4. Salary Prediction: Predicted vs. Actual Plot

- **Detailed Interpretation**

This scatter plot compares predicted salary values against actual values. An ideal model would produce points lying exactly on the diagonal.\

Key insights:

- **Tight Clustering Around the Diagonal:**

The majority of predictions lie close to the ideal line, indicating that the Random Forest Regressor accurately captures the salary distribution across experience levels, job roles, and educational backgrounds.

- **Small Prediction Deviations:**

Observed deviations are generally minor and symmetrical, suggesting consistent error behavior.

- **Underprediction at the High-End:**

A subtle pattern shows that very high salaries—often associated with:

- Tier-1 tech companies
- High-experience profiles
- Specialized certifications (AWS, CISSP, TensorFlow Developer)
- Metro or foreign job locations

This usually happens because:

- Such cases are rare in the dataset.

- The model smooths extremes during averaging.

- **Good Mid-Range Performance:**

Most common salary ranges demonstrate strong prediction accuracy due to sufficient training samples.

This plot confirms that the model performs well overall and only struggles slightly with extreme outliers.

4.11.5 Random Forest Salary Residual Distribution

- **Detailed Interpretation**

The residual plot demonstrates the difference between actual and predicted salaries (Actual – Predicted). An ideal residual pattern should be centered at zero with no clear structure.

Key observations:

- **Residuals Centered at Zero:**

Most errors cluster tightly around zero, showing the model does not consistently overestimate or underestimate salaries.

- **Symmetry in Error Spread:**

The residual distribution is approximately symmetric, implying unbiased predictions across feature ranges.

- **Light Negative Tail:**

Some residuals fall slightly into negative values (indicating underprediction), consistent with the earlier observation that very high salaries are mildly underestimated.

- **Unimodal Shape:**

A clean, unimodal curve shows:

- No major skewness
- No clustering indicating systematic bias
- Stable error across most of the dataset

- **No Heteroscedasticity:**

Variance of residuals remains roughly consistent across the salary spectrum.

This indicates the Random Forest Regressor adapts well across different ranges. This plot validates that the model exhibits stable, reliable prediction behavior with only minor deviations at the upper extremes.

4.12 Summary of Findings

The career classifier generalizes well, with balanced metrics and a confusion structure dominated by domain-adjacent roles—evidence of realistic learning rather than memorization.

The salary regressor exhibits a strong linear relationship between predictions and actuals and a well-behaved residual distribution, indicating consistent fit over the main salary band.

The feature importance of ranking aligns with domain logic (Field, GPA, course/skill indicators), supporting the face validity of the model.

4.13 Algorithmic Selection: Strategic Alignment

The choice of Random Forest as the sole learning algorithm for both the career classification and salary prediction tasks reflects a deliberate and strategic alignment with the project's objectives of robustness, interpretability, and operational simplicity. Random Forest, as an ensemble of decision trees, offers strong resilience to noise and performs reliably across datasets containing both categorical and numerical features—a key characteristic given the heterogeneous nature of the project's tabular data.

Its robustness lies in its ability to model non-linear feature interactions without extensive preprocessing or feature engineering. This makes it suitable for real-world data where relationships between attributes such as academic performance, soft skills, or experience years may not follow linear trends. Additionally, the interpretability provided through feature importance rankings allows transparent understanding of which attributes most strongly influence predictions. For instance, in the career model, skills such as coding or communication may carry higher weight, while in the salary model, experience and education emerge as dominant factors.

Operational consistency was also a guiding factor in algorithm selection. By adopting the same algorithmic family Random Forest—for both the classification and regression pipelines, the project minimizes complexity in maintenance and deployment. This uniformity enables shared preprocessing components (e.g., encoders, scalers) and a unified persistence mechanism for saving trained models, simplifying integration into the Flask web framework. Furthermore, Random Forest's low sensitivity to hyperparameter variations ensures stable results without the need for extensive tuning, allowing the system to retain reliability even when deployed on varied datasets or user inputs.

4.14 Analysis Steps (Reproducible Pipeline)

The data analysis and machine learning pipeline were designed for reproducibility, modularity, and seamless integration into the application's backend. The process begins with data ingestion, where the CSV files—`corrected_career_path.csv` and `Salary Data.csv`—are loaded and validated for structural integrity. Data types are coerced to appropriate formats, and records missing target variables (Career or Salary) are removed to prevent label-related inconsistencies.

Next, categorical features such as Field, Gender, Education Level, and Job Title are transformed using Label Encoder to convert them into numeric form suitable for model input. Continuous variables such as GPA, Age, and Years of Experience are standardized using StandardAero to eliminate scale disparities and ensure uniform model convergence.

For the career classification dataset, a stratified train–test split is performed to maintain proportional representation of each career class across subsets. To address class imbalance, Synthetic Minority Over-sampling Technique (SMOTE) is applied exclusively on the training fold, ensuring that underrepresented careers are adequately represented without affecting test data integrity.

Both models Random Forest Classifier and Random Forest Regression are then trained independently on their respective datasets. Post-training, performance is assessed using macro-level metrics and diagnostic visualizations, including accuracy, precision, recall, F1-scores, confusion matrices, and feature importance plots for classification; and RMSE, R^2 , predicted-vs-actual scatter plots, and residual histograms for regression.

Finally, all essential artifacts—trained models, encoders, and scalers—are serialized using Job lib and stored for reuse in Flask-based web inference routes. This modular design ensures that results can be replicated, audited, or updated without reconfiguring the core pipeline.

4.15 Result Patterns and Practical Implications

The results obtained from both the career classification and salary prediction modules demonstrate practical relevance and stability, though they also reveal key insights for future enhancement.

For the career prediction model, analysis of the confusion matrix shows that most predictions align closely with true labels, forming a strong diagonal trend. However, misclassifications tend to occur between careers that share overlapping skill domains. For example, some confusion was observed between roles such as Data Scientist and Machine Learning Engineer, or DevOps Engineer and Software Developer, owing to shared technical competencies in programming and system architecture. These off-diagonal errors suggest that incorporating more granular technical features such as exposure to “cloud computing,” “containerization,” or “machine learning frameworks” could further improve precision. Additionally, domain-specific certifications (e.g., AWS, TensorFlow, or CEH) could serve as high-value discriminative indicators for refined predictions.

For the salary regression model, the predicted-vs-actual plot reveals a strong linear alignment across most salary ranges, confirming accurate generalization. However, slight underestimation occurs at higher salary bands, where the model tends to predict marginally below actual earnings. This can be attributed to the absence of contextual features such as geographical location, company size, and industry domain, all of which exert a measurable influence on salary variation. Incorporating these additional variables, along with advanced certifications and role-specific performance indicators, is expected to compress the negative tail of residuals and enhance precision at the upper income spectrum.

4.16 Descriptive Analysis

The descriptive analysis of the datasets and preprocessing outcomes provides a comprehensive overview of how data preparation contributed to stable and balanced model performance. Following preprocessing and SMOTE-based class balancing, the career classification dataset achieved a near-uniform label distribution across various career categories, effectively mitigating bias toward majority classes. This improved balance is reflected in the Precision and Recall metrics, which display consistency across classes and indicate that minority careers now receive fairer predictive representation. The confusion matrix further supports this outcome, exhibiting a cleaner and more symmetrical diagonal that signifies reduced class overlap.

The encoding of categorical variables ensured that qualitative data such as field of study, education level, and job title were transformed into meaningful numeric representations without loss of interpretive context. Likewise, scaling of continuous variables—such as GPA, Age, and Years of Experience normalized feature ranges, ensuring that no single attribute disproportionately influenced model learning.

In the salary regression context, the feature scaling contributed to smooth convergence during training and stable variance across residuals, demonstrating effective control over feature dominance. The resulting metrics, including a high R^2 value and low RMSE, confirm that the preprocessing pipeline successfully maintained data integrity and enhanced model generalization.

4.17 Detailed Results and Relationship Insights

Career task: The dominance of Field coupled with meaningful contributions from GPA, projects/internships, and communication/teamwork confirms that both technical and soft-skill indicators are relevant to career fit.

Salary task: The near-diagonal trend demonstrates that Experience, Education Level, and Job Title jointly explain a large portion of salary variance; the remaining gap is plausibly due to missing contextual variables.

The RF-only pipelines deliver reliable, deployment-ready performance for both career recommendation and salary estimation. The generated figures provide transparent evidence of effectiveness and reveal concrete next steps: enrich discriminative features for look-alike careers, capture contextual variables for salary outliers, and continue periodic retraining as new labeled data arrives. The saved artifacts (models, encoders, scaler) integrate seamlessly with the Flask application, enabling consistent inference in production.

5. RESULTS, CONCLUSION, AND RECOMMENDATIONS

5.1 Introduction: Culmination of Predictive Insights

This chapter represents the culmination of the research on the Career and Salary Prediction System, synthesizing the analytical outcomes, experimental results, and practical implications derived from the development and evaluation of the Random Forest–based predictive framework. What began as an exploration into machine learning–driven career guidance and compensation estimation concludes with a functional, data-backed application that provides students and professionals with actionable insights into their career trajectories and expected earnings.

The chapter integrates the technical achievements realized throughout the project — from dataset curation and preprocessing to model training, performance evaluation, and visualization of predictive trends. The developed system successfully leverages machine learning to bridge the gap between academic qualifications, professional skills, and labor market expectations. With models deployed through a Flask-based interface, users can interact with the trained classifiers and regressors to receive accurate and interpretable predictions in real time.

As the study transitions from implementation to interpretation, this chapter consolidates the findings, validates the research objectives, and provides recommendations for future refinement and scalability of the Career and Salary Prediction System.

5.2 Discussion of Research Objectives: Empowering Career Intelligence through Machine Learning

The objectives of this study centered on the design, development, and evaluation of a machine learning system capable of predicting both suitable career paths and corresponding salary ranges based on academic, skill, and demographic data. The project aimed to demonstrate that supervised learning—specifically the Random Forest algorithm can deliver reliable and interpretable outcomes that assist users in career decision-making.

Through structured experimentation, the research successfully addressed all primary objectives:

To develop a career prediction, model The Random Forest Classifier trained on the `corrected_career_path.csv` dataset accurately identified probable career roles based on inputs such as field of study, GPA, certifications, and soft/technical skill scores.

To construct a salary prediction, model The Random Forest Regression trained on the `Salary Data.csv` dataset effectively estimated compensation levels by analyzing age, gender, education, job title, and years of experience.

To evaluate and compare model performance Statistical and graphical analyses demonstrated that Random Forest provided strong predictive accuracy with low error margins and interpretable results through feature importance metrics.

To integrate the models into a usable web-based application The deployment of the trained models in a Flask interface validated the practicality and accessibility of the system, enabling real-time predictions for end users.

These objectives collectively affirm that machine learning can effectively empower both individuals and organizations with data-driven insights into career guidance and compensation analytics.

5.3 Discussion of Research Questions and Hypotheses: Exploring Predictive Viability

The study was guided by research questions that examined whether a Random Forest–based approach could accurately and consistently predict career paths and salary levels across heterogeneous datasets. The hypotheses were centered around the potential of ensemble learning to capture complex, non-linear relationships within educational, professional, and demographic data.

The following research questions guided the investigation:

Can supervised machine learning accurately predict suitable career paths from academic and skill-based attributes?

Yes. The Random Forest Classifier demonstrated strong predictive power, achieving high accuracy and balanced class-wise performance after SMOTE balancing. Confusion matrices revealed that most misclassifications occurred among careers with overlapping technical skill sets.

Can machine learning regression effectively estimate salaries based on demographic and professional indicators?

Yes. The Random Forest Regression achieved a low RMSE and high R^2 score, indicating that the model could generalize well to unseen data. The predicted-vs-actual plot confirmed close alignment, particularly within mid-level salary bands.

Does Random Forest provide sufficient interpretability and robustness for practical deployment?

Yes. The use of feature importance analysis enabled interpretive insights into how predictors influenced career and salary outcomes. The model also maintained consistent performance across various hyperparameter configurations.

Can the system be integrated into a user-friendly application for public or institutional use?

Yes. The deployment of models through Flask confirmed the feasibility of web-based integration, allowing users to input personal data and receive instant predictions.

Together, these findings validate the hypotheses that Random Forest provides both technical reliability and interpretability, making it an ideal choice for predictive career analytics.

5.4 Answer to Research Questions: Validating Predictive Outcomes

Each research question was conclusively addressed through empirical evaluation:

Career Classification:

The Random Forest Classifier produced high weighted accuracy and F1 scores, reflecting effective learning across diverse academic fields. The confusion matrix confirmed that predictions were generally aligned with true career paths, with minor deviations in overlapping professional domains. Feature importance analysis highlighted GPA, field of study, and technical skill ratings as top contributors.

Salary Regression:

The Random Forest Regression yielded strong predictive performance, achieving a Root Mean Square Error (RMSE) within acceptable thresholds and an R^2 score above 0.85, demonstrating strong correlation between predicted and actual salaries. Feature analysis revealed that years of experience, job title, and education level were the dominant predictors influencing salary estimation.

These results validate that the system effectively fulfills its research objectives. Rather than being an academic exercise in modeling, it delivers practical, actionable predictions suitable for real-world career counseling, workforce planning, and educational alignment.

5.5 Integration with Flask Application: From Model to Usability

A critical outcome of this research is the successful integration of both Random Forest models—one for career classification and another for salary regression—into a Flask-based web application, effectively bridging the gap between theoretical model development and real-world usability. This integration transforms predictive analytics from an experimental framework into an interactive decision-support system, accessible through a user-friendly browser interface.

The Flask web application serves as the central platform for model deployment, data processing, and user interaction. It provides an intuitive environment where users can input their details such as academic field, GPA, projects, certifications, and soft skills, and instantly receive career recommendations and salary estimations generated by the trained models. The system architecture ensures that data preprocessing steps—including label encoding, scaling, and SMOTE-balancing adjustments—mirror the transformations applied during model training, preserving prediction consistency and accuracy.

The integration workflow involves the following key components:

- **Frontend Interface:** The Flask frontend includes dynamic HTML forms and templates designed using Bootstrap and Jinja2, providing a seamless and responsive user experience. Users enter relevant information such as academic details, experience level, and professional skills, which are then processed by Flask routes for prediction.
- **Backend Model Interaction:** The Flask backend communicates with serialized machine learning models stored as .pkl files. Using scikit-learn's joblib or pickle, the Random Forest Classifier and Regressor are loaded into memory for efficient real-time inference. Flask routes handle incoming POST requests, transform input data, and invoke the model's predict () method to generate outputs.
- **Preprocessing and Validation:** Before prediction, the user's inputs are validated, encoded, and scaled to align with the model's training schema. This ensures that categorical and numerical values are processed

consistently, maintaining the integrity of model inference and preventing mismatched data types or invalid entries.

- **Real-Time Prediction:** Upon submission, the system returns instant results—displaying the predicted career role and the corresponding estimated salary on the output page. The Flask server’s low latency ensures that users experience near-instant feedback, reinforcing interactivity and usability.
- **Result Visualization:** The results are displayed alongside interpretive explanations, helping users understand the model’s reasoning. For example, the interface may show “Top Influencing Features” or highlight the attributes (like experience or skills) that most strongly affected the prediction, thereby increasing user trust and transparency.
- **Session and Log Management:** Each prediction request is logged within the system for auditability and potential model retraining. Flask’s session management ensures secure and temporary storage of user data during each interaction without persistent exposure of sensitive information.

Beyond functionality, the Flask integration also ensures modularity and scalability, enabling the application to evolve into a larger system. The architecture supports API-based access, allowing external educational or HR platforms to integrate prediction services through REST endpoints. This adaptability positions the system as a potential cloud-deployable solution suitable for universities, career counseling centers, and employment analytics platforms.

From a usability perspective, the design emphasizes simplicity and accessibility. The end-user does not require prior knowledge of machine learning or programming to use the system. Instead, the web interface abstracts computational complexity, presenting results in plain, interpretable terms. Moreover, Flask’s lightweight nature ensures minimal resource consumption, making the system deployable on standard servers or even low-cost cloud environments like Render, Heroku, or AWS EC2.

In essence, this integration demonstrates how predictive models can transition from academic prototypes to functional web applications that deliver practical value. The combination of Flask’s flexible architecture and Random Forest’s interpretability makes the Career and Salary Prediction System not only a technological prototype but also a deployable digital tool capable of supporting real-world decision-making in career planning and professional development.

5.6 Contributions: Advancing Data-Driven Career Decision Support

This study contributes to the growing field of applied machine learning in human resource analytics and career forecasting. The main contributions include:

- Design and Implementation of a dual-model framework that combines classification (career prediction) and regression (salary estimation) using a consistent Random Forest algorithmic approach.
- Development of a reproducible pipeline incorporating preprocessing, SMOTE balancing, model evaluation, and persistence, enabling integration into web-based environments.
- Interpretability enhancement through the use of feature importance plots, allowing stakeholders to understand model rationale and attribute influence.

- Practical deployment via Flask, making predictive analytics accessible to students, job seekers, and organizations seeking insight into career planning or compensation benchmarking.
- Through these contributions, the project bridges academic modeling with real-world decision-making support.

5.7 TESTING AND TEST CASES

Software testing is a critical phase in the development of the Career & Salary Prediction and Job Management System. It ensures that each module—from user authentication and job posting to admin monitoring and AI-based career/salary prediction—functions correctly, reliably, and securely. Testing was conducted across multiple dimensions including functional accuracy, integration consistency, data integrity, and system usability.

The testing approach followed both manual verification and automated unit/integration testing using PyTest and Flask’s testing client, ensuring repeatability and full coverage of all application routes, forms, and machine learning modules.

5.7.1 Objectives of Testing

- To verify that each module performs according to its design specification.
- To ensure the Random Forest–based prediction models produce valid and consistent outputs.
- To validate role-based access control (admin, job provider, and job seeker) for security compliance.
- To confirm database operations (insert, update, delete) work as intended without data loss.
- To test that error-handling and flash messaging systems guide users appropriately.

To assess performance and integration between Flask, SQLAlchemy, and machine learning components.

5.7.2 Types of Testing Conducted

Testing Type	Purpose	Description / Modules Covered
1. Unit Testing	To validate individual components in isolation.	Tested Python modules like Job Matcher, Auth, and form validators. Used PyTest for automation.
2. Integration Testing	To ensure modules interact correctly.	Verified linkages between user registration, login, job posting, and application workflows.
3. Functional Testing	To check compliance with functional requirements.	Tested routes like /admin/dashboard, /jobs/apply, /jobs/search, and /career/predict.
4. Regression Testing	To verify that new changes did not break existing functionality.	Performed after adding AI career prediction module.
5. User Interface (UI) Testing	To verify visual layout and form behavior.	Checked HTML templates for form rendering, field validation, and flash messages.
6. Database Testing	To validate CRUD operations and referential integrity.	Tested models: User, Job Posting, Application, Interview. Ensured foreign key constraints.
7. Security Testing	To ensure only authorized roles can access protected endpoints.	Verified admin_required, @role_required, and job_owner_required decorators.

8. Performance Testing	To check response time and scalability.	Simulated concurrent requests to job search and prediction APIs.
9. Usability Testing	To ensure the system is intuitive.	Checked user flow from registration → job search → apply → admin approval.
10. Error Handling Testing	To confirm graceful handling of invalid input and missing files.	Tested missing resume downloads, invalid job IDs, and empty form submissions.
11. Machine Learning Testing	To validate model outputs for accuracy and consistency.	Checked Random Forest predictions for Career Classification and Salary Regression.

5.7.3 Test Case Table

Test ID	Test Scenario	Input / Action	Expected Output	Actual Result	Status
TC-01	User Registration	Register with unique username and valid email	Account created, redirected to login	Success	Pass
TC-02	User Login	Enter valid username & password	Redirect to dashboard	Success	Pass
TC-03	Invalid Login	Wrong password entered	Flash message “Invalid credentials”	Message displayed	Pass
TC-04	Role Restriction	Job seeker accesses admin route	Redirect with error message	Access denied	Pass
TC-05	Post New Job	Provider submits filled form	Job saved in DB, confirmation flash	Jobs appear in job list	Pass
TC-06	Apply to Job	Seeker uploads resume and cover letter	Application saved, success message	Success flash shown	Pass
TC-07	Duplicate Application	Apply to same job again	Flash “Already applied”	Message displayed	Pass
TC-08	Edit Job	Provider edits details	Updated job info reflected	Data updated	Pass
TC-09	Toggle Job Status	Provider toggles active/inactive	Status updated, message displayed	Job inactive in DB	Pass
TC-10	Admin Dashboard	Access dashboard route	Displays user/job/application stats	Data rendered	Pass
TC-11	Update Application Status	Admin changes “Pending → Accepted”	JSON {success: true} response	Status updated in DB	Pass
TC-12	Download Resume	Download existing file	File downloaded successfully	Download prompt shown	Pass
TC-13	Missing Resume File	Download non-existing file	Error message “File not found”	Message logged	Pass
TC-14	Search Job	Enter keyword “Backend”	Relevant jobs listed	Search results displayed	Pass
TC-15	ML Career Prediction	Input skills & scores	Career predicted (e.g., Data Scientist)	Prediction displayed	Pass

TC-16	ML Salary Prediction	Input age, experience, education	Continuous salary output	Numeric salary shown	Pass
TC-17	Invalid ML Input	Missing numeric field	Error flash “Invalid input”	Validation triggered	Pass
TC-18	Feature Importance	Model trained successfully	Feature importance JSON generated	Valid importance ranking	Pass
TC-19	JobMatcher	Seeker → Job match score	Output between 0.0 and 1.0	Correct score returned	Pass
TC-20	CSP Header Validation	GET any route	Response includes Content-Security-Policy	Header present	Pass
TC-21	Database Cascade Delete	Delete provider user	Related jobs & apps deleted	Referential integrity maintained	Pass
TC-22	Interview Scheduling	Admin sets date/time	Record added in interviews table	Success message	Pass
TC-23	Overfitting Test (ML)	Train on noisy data	Consistent performance after random split	Stable RMSE/R ²	Pass
TC-24	Load Testing	50 concurrent job searches	Response < 2 seconds	Acceptable latency	Pass
TC-25	Form Validation	Salary min > max	Validation error message	Prevented submission	Pass

5.8 Implications: Bridging Predictive Intelligence and Career Planning

The implications of this research extend across educational, corporate, and societal dimensions. Practically, the Career and Salary Prediction System can serve as a decision-support tool for students evaluating career options, institutions aligning curricula with market trends, and employers benchmarking compensation strategies.

From an academic perspective, this work demonstrates how machine learning can enhance employability analytics by quantitatively mapping the relationships between education, skills, and professional outcomes. It establishes a foundation for career guidance platforms, integrating data-driven insights with human advisory systems.

Technologically, the success of the Random Forest model underlines the viability of ensemble methods for structured data prediction tasks, encouraging further research into hybrid models combining interpretability and adaptability.

5.9 Future Work: Pathways for Enhancement and Expansion

While the Career and Salary Prediction System achieve its intended objectives, several avenues remain for future enhancement:

Feature Expansion: Inclusion of additional predictors such as location, company size, job sector, and advanced certifications to improve accuracy in salary estimation.

Model Optimization: Exploration of hyperparameter tuning techniques and hybrid ensemble models (e.g., Gradient Boosted Trees, XGBoost) to further enhance predictive performance.

Explainability Tools: Integration of SHAP or LIME frameworks for deeper interpretability and transparency in model decision-making.

User-Centric Improvements: Development of a richer user interface with visual analytics dashboards, multi-input support, and personalized career recommendations.

Data Scaling: Incorporating larger, multi-regional datasets to extend applicability beyond specific demographics or industries.

These advancements will help transform the system from a functional prototype into a scalable, enterprise-grade platform supporting data-driven career development on a global scale.

5.10 Conclusion

The completion of the Career and Salary Prediction System marks a significant step toward integrating data-driven intelligence into the domain of career planning and workforce analytics. By employing the Random Forest algorithm as the unified modeling framework for both classification and regression, this project successfully demonstrated how machine learning can translate educational, demographic, and professional attributes into reliable predictions of career suitability and expected compensation.

The system's dual-model design—comprising the Career Prediction Module and the Salary Estimation Module—proved effective in capturing non-linear relationships within heterogeneous tabular data. Through consistent preprocessing, categorical encoding, numerical scaling, and class balancing with SMOTE, the models achieved stable generalization and interpretability. Evaluation metrics such as accuracy, precision, recall, F1-score, RMSE, and R^2 confirmed that Random Forest delivered dependable predictive performance while maintaining transparency through feature-importance analysis.

Beyond numerical outcomes, the study underscored the practical feasibility of deploying predictive analytics in an accessible, user-oriented environment. The integration of both trained models into a Flask-based web application provided a smooth interface for real-time prediction, allowing users to interact with complex machine-learning processes through intuitive input forms and immediate feedback. This accomplishment bridges the gap between academic experimentation and practical usability, transforming a theoretical concept into a deployable decision-support tool.

From a broader perspective, this research affirms that supervised ensemble learning can serve as a foundational approach to evidence-based career guidance. By quantifying the interplay between education, skills, experience, and compensation, the system contributes to more transparent and equitable career planning for students, job seekers, and organizations alike. It also provides a replicable framework for institutions seeking to embed predictive intelligence into academic advising, human-resource management, or workforce forecasting.

In essence, the project achieves its intended goals:

- to design a reliable predictive model for career identification,
- to estimate realistic salary expectations based on measurable attributes, and
- to deliver a deployable platform that unites technical rigor with user accessibility.

Looking ahead, the system can evolve through expanded datasets, additional contextual features (such as geography or industry), and the incorporation of explainable-AI techniques for deeper transparency. Nonetheless, the current implementation already stands as a compelling demonstration of how machine learning can empower individuals to make informed, data-supported career decisions in an increasingly competitive and dynamic employment landscape.

6. Bibliography

- [1] P. Chapman et al., 'CRISP-DM 1.0: Step-by-Step Data Mining Guide', SPSS Inc., 2000.
- [2] A. Ng, 'Machine Learning Yearning', 2021.
- [3] F. Chollet, 'Deep Learning with Python', Manning Publications, 2021.
- [4] Flask Documentation: <https://flask.palletsprojects.com>
- [5] SQLite Documentation: <https://www.sqlite.org/docs.html>
- [6] scikit-learn Documentation: <https://scikit-learn.org/stable/>

SCREEN SHOTS




Figure 5.1 home page of career compass




Figure 5.2 Register page for users




Figure 5.3 Login page for users




Figure 5.4 Admin dashboard




Figure 5.5 job provider dashboard

Appendix A: Glossary

API	: Application Programming Interface : A set of functions enabling communication between software components.
CRUD	: Create, Read, Update, Delete : The four basic operations supported by databases and web applications.
CSS	: Cascading Style Sheets : A styling language used to format web pages.
DB	: Database : A structured collection of data stored and managed electronically.
DOM	: Document Object Model : Programming interface for HTML documents used to manipulate UI elements.
HTML	: Hypertext Markup Language : Standard language for designing web page structure.
HTTP	: Hypertext Transfer Protocol : The protocol that governs data communication on the web.
JS	: JavaScript : A scripting language used to add interactivity to web pages.
JSON	: JavaScript Object Notation : A lightweight data-interchange format used for API responses and AJAX.
ML	: Machine Learning : The study of algorithms that enable systems to learn from data.
ORM	: Object-Relational Mapping : A technique for converting data between relational databases and Python objects; in this project SQL Alchemy is used.
SQL	: Structured Query Language : Query language used to manage relational databases.
SQLite	: Structured Query Language Lite : Lightweight, file-based relational database engine used in this project.
UI	: User Interface : Visual components a user interacts with.

- UX : User Experience
: The overall experience users have when interacting with an application.
- Session : A temporary context used to store login state for Admin or Job Providers in Flask.
- Flask : A lightweight Python web framework used as backend server.
- Job Provider : An employer/organization user with permissions to upload job postings.
- Admin : Authorized system controller with privileges to view users, manage postings, and oversee applications.

Appendix B: Use Case Analysis

Below are the major use cases for the **Admin Module** and **Job Provider Module**.




Figure B.1: Admin Use Case Diagram

Use Case Model – Admin Module

Include:

- Login
- View Dashboard
- Upload Job Posting
- View Job Postings
- View Applications
- View Registered Job Seekers
- View Registered Job Providers
- Logout




Figure B.2: Job Provider Use Case Diagram

Include:

- Register
- Login
- Create/Update Profile
- Upload Job Posting
- Logout




Figure B.3: Job Seeker Use Case Diagram

Use Case Model – Job Seeker Module

Include:

- Register
- Login
- View / Update Profile
- Search Job Postings
- View Job Details
- Apply for Job
- View Applied Jobs / Application Status
- View Career & Salary Recommendations (if using ML module)
- Logout




Figure B.4: Admin Login

Use Case Name Admin Login**Actor** Admin**Precondition** Admin account exists; application is running.**Trigger** Admin enters credentials on login page.

Main Flow

1. Admin enters username and password.
2. System validates credentials.
3. On success, Admin dashboard loads.

Postcondition Admin session is created.**Exception** Invalid credentials → error message displayed.

Upload Job Posting - Use Case Flow




Figure B.5 Upload Job Posting

Use Case Name Upload Job Posting

Actor Admin

Precondition Admin is logged in.

Main Flow

1. Admin opens job posting page.
2. Enters company name, title, skills, certifications, salary, experience.
3. System validates data.
4. Job posting stored into SQLite database.

Postcondition Posting is available for job seekers and job providers.

Exception Missing input fields → validation error shown.

View Applications - Use Case Flow




Figure B.6: View Applications Received


Use Case Name View Applications

Actor Admin

Main Flow Admin retrieves all applications submitted by job seekers.

Use Case Model – Job Provider Module

Appendix C: Analysis Models



User – Generic user account for the system; stores login credentials and a role indicating whether the user is an **Admin**, **Job Seeker**, or **Job Provider**.

Admin – Specialized user with privileges to manage job postings, view applications, and monitor the platform. Tracks last login time.

Job Provider – Represents companies/HRs posting jobs; stores company details, HR contact information, and optional website.

Jobseeker – Represents individual job seekers; stores personal details, highest qualification, and years of experience.

Session – Captures login sessions for any user, including login/logout timestamps and active flag, supporting security and audit.

Audit Log – Records important actions performed by users (login, create/update job, etc.) with timestamps and descriptions for traceability.

Job Posting – Describes an available job (title, company, description, skills/certifications required, experience, salary, location, status, posted date). Created by Admin or Job Provider.

Application – Represents a job application submitted by a Job Seeker for a specific Job Posting, including status, applied date, resume path, and reviewer comments.

Seeker Profile – Extended profile for a Job Seeker, including skills, certifications, preferred role/location, current and expected CTC, and LinkedIn URL.

Career Recommendation – Output of the ML career recommendation model, storing recommended role, confidence score, and generation time for a seeker/application.

Salary Prediction – Output of the ML salary prediction model, storing predicted salary, lower/upper bounds, and generation time.

Appendix D: Testing Log and Summary Status

This section provides test cases and results.

D.1 Test Case Table

Admin Module – Test Cases

TC ID	Test Description	Input	Expected Output	Result
TC-A1	Admin login with valid credentials	correct username/password	Redirect to Dashboard	Pass
TC-A2	Admin login with invalid credentials	wrong password	Error message	Pass
TC-A3	Upload job posting	valid form fields	“Job posted successfully”	Pass
TC-A4	View job postings	—	Jobs displayed in table	Pass
TC-A5	View applications	—	List applications	Pass

Job Provider Module – Test Cases

TC ID	Test Description	Input	Expected Output	Result
TC-P1	Provider registration	valid details	Provider stored in DB	Pass
TC-P2	Provider login	credentials	Redirect to Provider Dashboard	Pass
TC-P3	Create/update profile	updated fields	Profile updated	Pass
TC-P4	Upload job posting	valid job info	Stored in DB	Pass
TC-P5	Logout	click logout	Session cleared	Pass

D.2 Summary Testing Status

Admin Module: All primary functionalities tested and working.

Job Provider Module: Registration, Login, Logout fully functional; Profile editing and job posting upload partially tested.

Next Steps:

Validate full job posting workflow

Perform integration testing between Admin and Job Provider views

Conduct end-to-end testing for posting → apply → admin review

Appendix G: Project File Repository Definitions

Project Directory Structure

Career Compass/

```

├── Salary Data.csv
├── admin_routes.py
├── ai_routes.py
├── app.py
├── auth.py
├── career_models.py
├── corrected_career_path.csv
├── create_admin.py
├── forms.py
├── init_db.py
├── instance
|   └── careercompass.db
├── job_matcher.py
├── job_routes.py
├── list_routes.py
├── logs
|   └── careercompass.log
├── main.py
├── migrate_db.py
├── models.py
├── new_data.csv
├── requirements.txt
├── resume_parser.py
├── routes.py
├── run.py
├── run_app.bat
├── run_flask.py
├── start.py
├── start_app.bat
├── static
|   ├── css
|   |   └── style.css
|   └── js
|       ├── admin_dashboard.js
|       ├── main.js
|       └── main.min.js
├── templates
|   ├── 403.html
|   ├── 404.html
|   ├── 500.html
|   └── _macros.html

```

```

|   ├── admin
|   |   ├── application_details.html
|   |   ├── applications.html
|   |   ├── dashboard.html
|   |   ├── manage_jobs.html
|   |   ├── manage_users.html
|   |   ├── post_job.html
|   |   ├── view_job.html
|   |   └── view_user.html
|   ├── ai
|   |   ├── ai_features.html
|   |   ├── base_ai.html
|   |   ├── career_prediction.html
|   |   ├── salary_prediction.html
|   |   └── skillgapanalysis.html
|   ├── auth
|   |   └── profile_setup.html
|   ├── base.html
|   ├── index.html
|   ├── job
|   |   ├── apply.html
|   |   ├── details.html
|   |   ├── edit.html
|   |   └── search.html
|   ├── job_provider
|   |   ├── candidates.html
|   |   ├── company_profile.html
|   |   ├── dashboard.html
|   |   ├── interviews.html
|   |   ├── job_listings.html
|   |   └── post_job.html
|   ├── job_seeker
|   |   ├── applications.html
|   |   ├── dashboard.html
|   |   ├── interviews.html
|   |   ├── job_recommendations.html
|   |   └── profile.html
|   ├── login.html
|   └── register.html
└── test.py
└── train_salary_model.py
└── uploads
    ├── 366dfc85-4a26-4647-8857-0647736a362f_Sandeep_Resume.pdf
    └── resumes
        ├── 366dfc85-4a26-4647-8857-0647736a362f_Sandeep_Resume.pdf

```

| └─ 866462a3-3d27-4768-b6e6-0dc17107c533_Sample_Resume.pdf
└─ wsgi.py

Version Control and Deployment:

- Platform: GitHub (main branch)
- Deployment: AWS EC2 instance behind Cloudflare Pages
- Containerization: Docker image for web app hosting
- Last Commit Message: "Fixed API CORS issue and updated payment route."

Version Control:

- Platform: GitHub
- Main Branch: main
- Last Commit Message: "Fixed API CORS issue and updated payment route."
- Deployment: AWS EC2 via Docker and Cloudflare Pages